

PRACTICAL - 1

Aim: Implement code to

- i. read an image, observe image as function along with its attributes, create a sample image and store it for grayscale images
- ii. read an image, observe image as function along with its attributes, split image in different color image planes, merge splitted image plane and observe effect on merged image by changing red, green and blue image plane values
- iii. convert color image to grayscale and HSV image and observe HSV values
- iv. create images of different shapes
- v. resize image with different types of interpolation

Program:

i.

% Read an image and observe its attributes

imagePath = 'saulgoodman.png';

originalImage = imread(imagePath);

disp(['Original Image Size: ', num2str(size(originalImage))]);

disp(['Image Data Type: ', class(originalImage)]);

% Create a sample grayscale image

sampleImage = randi([0, 255], [300, 300], 'uint8');

imwrite(sampleImage, 'sample_grayscale_image.jpg');

Output:

```
Original Image Size: 393 439 3
```

```
Image Data Type: uint8
```

Program:

ii.

% Read a color image and observe its attributes

```
colorImagePath = 'saulgoodman.png';
```

```
colorImage = imread(colorImagePath);
```

```
disp(['Color Image Size: ', num2str(size(colorImage))]);
```

```
disp(['Color Image Data Type: ', class(colorImage)]);
```

% Split image into different color planes (R, G, B)

```
r = colorImage(:, :, 1);
```

```
g = colorImage(:, :, 2);
```

```
b = colorImage(:, :, 3);
```

% Merge the color planes

```
mergedImage = cat(3, r, g, b);
```

% Observe the effect by changing red, green, and blue image plane values

% For example, change red plane values:

```
r = r * 1.5;
```

% Merge the modified planes

```
modifiedImage = cat(3, r, g, b);
```

Output:

```
Color Image Size: 393  439    3
```

```
Color Image Data Type: uint8
```

Program:

iii.

% Convert color image to grayscale

grayImage = rgb2gray(colorImage);

% Convert color image to HSV

hsvImage = rgb2hsv(colorImage);

% Observe HSV values

disp(['Hue Values (min, max): ', num2str([min(hsvImage(:, :, 1)), max(hsvImage(:, :, 1))]))];

disp(['Saturation Values (min, max): ', num2str([min(hsvImage(:, :, 2)), max(hsvImage(:, :, 2))]))];

disp(['Value (Brightness) Values (min, max): ', num2str([min(hsvImage(:, :, 3)), max(hsvImage(:, :, 3))]))];

Output:

Hue Values (min, max): 0	0	0	0	0
Saturation Values (min, max): 0	0	0	0	0
Value (Brightness) Values (min, max): 0.054902	0.054902	0.054902	0.054902	0.054902

Program:



iv.

% Create images of different shapes

rectangleImage = zeros(300, 400);

rectangleImage(50:250, 100:300) = 255;

```
diamondImage = zeros(300, 400);  
[x, y] = meshgrid(1:400, 1:300);  
diamondMask = abs((x - 200) / 200) + abs((y - 150) / 150) <= 1;  
diamondImage(diamondMask) = 255;  
% Display or save the created images  
imwrite(rectangleImage, 'rectangle_image.jpg');  
imwrite(diamondImage, 'diamond_image.jpg');
```

Output:

Program:

v.

% Read an image to resize

imagePathToResize = 'saugoodman.png';

imageToResize = imread(imagePathToResize);

% Resize image with different types of interpolation

resizedNearest = imresize(imageToResize, 0.5, 'nearest');

resizedBilinear = imresize(imageToResize, 0.5, 'bilinear');

resizedBicubic = imresize(imageToResize, 0.5, 'bicubic');


% Save resized images


imwrite(resizedNearest, 'resized_nearest.jpg');


imwrite(resizedBilinear, 'resized_bilinear.jpg');

imwrite(resizedBicubic, 'resized_bicubic.jpg');

Output:

Name	Size	Class	Value
resized_bicubic	197×220×3	uint8	197×220×3 uint8
Preview:			
			

Name	Size	Class	Value
resized_nearest	197×220×3	uint8	197×220×3 uint8
Preview:			
			

Name	Size	Class	Value
resized_bilinear	197×220×3	uint8	197×220×3 uint8
Preview:			
			

Conclusion: In this practical, I performed image attribute observation, color channel manipulation, color space conversion, shape generation, and image resizing with different interpolation methods using OpenCV.

PRACTICAL - 2

Aim: Implement code to

- i. obtain negative of given images
- ii. perform contrast stretching on given images
- iii. perform thresholding on given images
- iv. perform Otsu's thresholding on given images
- v. perform log transformation on given images
- vi. perform gamma correction on given images

Program:

i.

% Read the original image

```
originalImage = imread('pexels-brett-sayles-6424244.jpg');
```

% Obtain negative of the image

```
negativeImage = 255 - originalImage;
```

% Display or save the negative image

```
imshow(negativeImage);
```

```
imwrite(negativeImage, 'negative_image.jpg');
```

Output:



Program:

ii.

```
% Read the original image
```

```
originalImage = imread('pexels-brett-sayles-6424244.jpg');
```

```
% Convert the original image to double data type
```

```
originalImage = double(originalImage);
```

```
% Perform contrast stretching
```

```
minIntensity = min(originalImage(:));
```

```
maxIntensity = max(originalImage(:));
```

```
stretchedImage = uint8((originalImage - minIntensity) * (255 / (maxIntensity -  
minIntensity)));
```

```
% Display or save the contrast-stretched image
```

```
imshow(stretchedImage);
```

```
imwrite(stretchedImage, 'contrast_stretched_image.jpg');
```

Output:

Program:

iii.

% Read the grayscale image

grayImage = imread('grayscale_image.jpg');

% Set a threshold value

threshold = 128;

% Perform thresholding

thresholdedImage = grayImage > threshold;

% Display or save the thresholded image

imshow(thresholdedImage);

imwrite(thresholdedImage, 'thresholded_image.jpg');

Output:

Program:

iv.

% Read the grayscale image

```
grayImage = imread('grayscale_image.jpg');
```

% Perform Otsu's thresholding

```
threshold = graythresh(grayImage);
```

```
otsuThresholdedImage = imbinarize(grayImage, threshold);
```

% Display or save the Otsu's thresholded image

```
imshow(otsuThresholdedImage);
```

```
imwrite(otsuThresholdedImage, 'otsu_thresholded_image.jpg');
```

Output:

Program:

v.

% Read the grayscale image

```
grayImage = imread('grayscale_image.jpg');
```

% Perform log transformation

```
c = 1; % Scaling factor
```

```
logTransformedImage = c * log(1 + double(grayImage));
```

% Scale the log-transformed image to [0, 255]

```
logTransformedImage = uint8((logTransformedImage /  
max(logTransformedImage(:))) * 255);
```

% Display or save the log-transformed image

```
imshow(logTransformedImage);
```

```
imwrite(logTransformedImage, 'log_transformed_image.jpg');
```

Output:

Program:

vi.

% Read the grayscale image

```
grayImage = imread('grayscale_image.jpg');
```

% Set gamma value (adjust as needed)

```
gamma = 1.5;
```

% Perform gamma correction

```
gammaCorrectedImage = imadjust(grayImage, [], [], gamma);
```

% Display or save the gamma-corrected image

```
imshow(gammaCorrectedImage);
```

```
imwrite(gammaCorrectedImage, 'gamma_corrected_image.jpg');
```

Output:

Conclusion: In this practical, I implemented various image manipulation techniques, including obtaining negative images, enhancing contrast through stretching, segmenting objects using thresholding, applying automatic thresholding with Otsu's method, adjusting contrast with logarithmic transformation, and fine-tuning brightness and contrast using gamma correction.

PRACTICAL - 3

Aim: Implement code to

- i. apply intensity slicing on given images
- ii. apply bit plane slicing on given images and observe information on different bit planes
- iii. calculate histograms on different contrast images.
- iv. apply normal histogram equalization and CLAHE histogram equalization on given images.
- v. apply histogram matching on given images

Program:

i.

% Read the original grayscale image

```
originalImage = imread('pexels-brett-sayles-6424244.jpg');
```

% Define lower and upper thresholds for intensity slicing

```
lowerThreshold = 100;
```

```
upperThreshold = 200;
```

% Perform intensity slicing

```
outputImage = originalImage;
```

```
outputImage(originalImage >= lowerThreshold & originalImage <= upperThreshold) = 255;
```

% Display or save the intensity-sliced image

```
imshow(outputImage);
```

```
imwrite(outputImage, 'intensity_sliced_image.jpg');
```

Output:**Program:**

ii.

% Read the original grayscale image

originalImage = imread('rectangle_image.jpg');

% Convert the image to binary representation

binaryImage = imbinarize(originalImage);

% Perform bit plane slicing

bitPlanes = zeros(size(binaryImage, 1), size(binaryImage, 2), 8);

for i = 1:8

bitPlanes(:, :, i) = bitget(originalImage, i);

end


% Display or save bit planes

for i = 1:8

subplot(2, 4, i);

imshow(bitPlanes(:, :, i) * 255);

```
title(['Bit Plane ', num2str(i)]);  
end
```

Output:**Program:**

iii.

```
% Read different contrast images
```

```
image1 = imread('download (1).jpeg');
```

```
image2 = imread('pexels-brett-sayles-6424244.jpg');
```

```
% Calculate histograms
```

```
histogramImage1 = imhist(image1);
```

```
histogramImage2 = imhist(image2);
```

```
% Display histograms
```


```
figure;
```

```
subplot(2, 1, 1);
```

```
bar(histogramImage1);
```

```
title('Histogram of Image 1');
```

```
subplot(2, 1, 2);  
bar(histogramImage2);  
title('Histogram of Image 2');
```

Output:**Program:**

iv.

```
% Read the original image
```

```
originalImage = imread('jesse.png');
```

```
% Convert the image to grayscale if it's a color image
```

```
if size(originalImage, 3) == 3
```

```
    originalImage = rgb2gray(originalImage);
```

```
end
```



```
% Perform normal histogram equalization
histeqImage = histeq(originalImage);


% Perform CLAHE (Contrast Limited Adaptive Histogram Equalization)
claheImage = adapthisteq(originalImage, 'ClipLimit', 0.02, 'Distribution',
'rayleigh');

% Display or save the processed images
imshowpair(histeqImage, claheImage, 'montage');
imwrite(histeqImage, 'histeq_image.jpg');
imwrite(claheImage, 'clahe_image.jpg');
```

Output:**Program:**

```
v.
% Read the reference image and the image to be matched
referenceImage = imread('download (1).jpeg');
imageToMatch = imread('jesse.png');
```

```
% Perform histogram matching  
matchedImage = imhistmatch(imageToMatch, referenceImage);  
  
% Display or save the matched image  
imshowpair(referenceImage, matchedImage, 'montage');  
imwrite(matchedImage, 'matched_image.jpg');
```

Output:

Conclusion: In this practical, I learnt image enhancement techniques, such as intensity slicing, bit plane slicing, histogram analysis, histogram equalization (standard and CLAHE), and histogram matching.

PRACTICAL - 4

Aim: Implement code to

- i. perform cross-correlation and convolution on images in spatial domain
- ii. apply smoothing spatial filters of different kernel sizes on images
- iii. apply sharpening spatial filters of different kernel sizes on images
- iv. apply non-linear spatial filters of different kernel sizes on images
- v. analyze noise removal with different smoothing spatial filters and non-linear filters
- vi. perform unsharp masking and high-boost filtering on different images

Program:

i.

% Read the original grayscale image and kernel

originalImage = imread('pexels-brett-sayles-6424244.jpg');

grayImage = rgb2gray(originalImage);

kernel = fspecial('gaussian', [3 3], 1); % Example Gaussian kernel

% Perform cross-correlation

crossCorrelationResult = xcorr2(double(grayImage), kernel);

% Perform convolution

convolutionResult = conv2(double(grayImage), kernel, 'same');

% Display or save the results

imshow(crossCorrelationResult, []);

imwrite(uint8(convolutionResult), 'convolution_result.jpg');

Output:**Program:**

ii.

% Apply smoothing spatial filters with different kernel sizes

```
smoothedImages = cell(1, 3);
```

```
kernelSizes = [3, 5, 7];
```

```
for i = 1:length(kernelSizes)
```

```
    kernel = fspecial('average', kernelSizes(i));
```

```
    smoothedImages{i} = conv2(double(grayImage), kernel, 'same');
```

```
end
```

% Display or save the smoothed images

```
for i = 1:length(kernelSizes)
```

```
    imshow(uint8(smoothedImages{i}), []);
```

```
    imwrite(uint8(smoothedImages{i}), ['smoothed_image_kernel_',  
num2str(kernelSizes(i)), '.jpg']);
```

```
end
```

Output:**Program:**

iii.

% Apply sharpening spatial filters with different kernel sizes

```
sharpenedImages = cell(1, 3);
```

```
kernelSizes = [3, 5, 7];
```

```
for i = 1:length(kernelSizes)
```

```
    % Create the kernel for sharpening
```

```
    kernel = -fspecial('average', kernelSizes(i)) + 2 * fspecial('gaussian',  
kernelSizes(i), 1);
```

```
    sharpenedImages{i} = conv2(double(grayImage), kernel, 'same');
```

```
end
```

```
% Display or save the sharpened images
```

```
for i = 1:length(kernelSizes)
```

```
    imshow(uint8(sharpenedImages{i}), []);
```

```
    imwrite(uint8(sharpenedImages{i}), ['sharpened_image_kernel_',  
num2str(kernelSizes(i)), '.jpg']);
```

```
end
```

Output:**Program:**

iv.

% Apply non-linear spatial filters with different kernel sizes

```
filteredImages = cell(1, 3);
```

```
for i = 1:length(kernelSizes)
```

```
    filteredImages{i} = medfilt2(grayImage, [kernelSizes(i), kernelSizes(i)]);
```

```
end
```

% Display or save the filtered images

```
for i = 1:length(kernelSizes)
```

```
    imshow(uint8(filteredImages{i}), []);
```

```
    imwrite(uint8(filteredImages{i}), ['filtered_image_kernel_',  
num2str(kernelSizes(i)), '.jpg']);
```

```
end
```

Output:**Program:**

v.

% Read the original image

```
originalImage = imread('jesse.png');
```

% Convert the original image to grayscale

```
grayOriginalImage = rgb2gray(originalImage);
```

% Add noise to the grayscale image

```
noisyImage = imnoise(grayOriginalImage, 'gaussian', 0, 0.01); % Adding Gaussian noise
```

% Apply smoothing spatial filters and non-linear filters to noisy image

```
smoothedImages = cell(1, length(kernelSizes));
```

```
filteredImages = cell(1, length(kernelSizes));
```

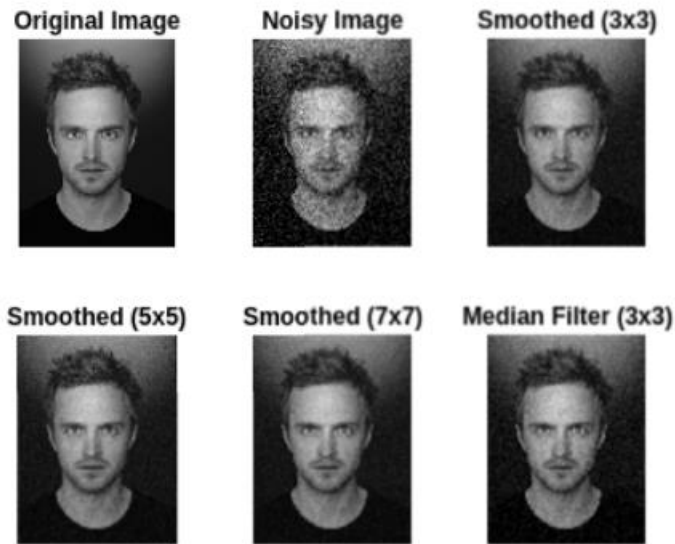
```
for i = 1:length(kernelSizes)
```

```
% Apply smoothing spatial filter (Gaussian blur)
gaussianKernel = fspecial('gaussian', kernelSizes(i), 1);
smoothedImages{i} = imfilter(noisyImage, gaussianKernel, 'replicate');

% Apply non-linear filter (median filter)
filteredImages{i} = medfilt2(noisyImage, [kernelSizes(i), kernelSizes(i)]);
end

% Display the results
figure;
subplot(2, 3, 1), imshow(grayOriginalImage), title('Original Image');
subplot(2, 3, 2), imshow(noisyImage), title('Noisy Image');
subplot(2, 3, 3), imshow(smoothedImages{1}), title('Smoothed (3x3)');
subplot(2, 3, 4), imshow(smoothedImages{2}), title('Smoothed (5x5)');
subplot(2, 3, 5), imshow(smoothedImages{3}), title('Smoothed (7x7)');
subplot(2, 3, 6), imshow(filteredImages{1}), title('Median Filter (3x3)');

% Save the processed images
imwrite(smoothedImages{1}, 'smoothed_3x3.jpg');
imwrite(smoothedImages{2}, 'smoothed_5x5.jpg');
imwrite(smoothedImages{3}, 'smoothed_7x7.jpg');
imwrite(filteredImages{1}, 'median_filter_3x3.jpg');
```


Output:**Program:**

vi.

% Read the original grayscale image

```
originalImage = imread('jesse.png');
```

% Apply unsharp masking for sharpening

```
smoothedImage = imgaussfilt(originalImage, 2); % Example Gaussian filter
```

```
unsharpMask = double(originalImage) - double(smoothedImage);
```

```
unsharpMaskedImage = double(originalImage) + 1.5 * unsharpMask;
```

% Apply high-boost filtering for sharpening

```
highBoostMaskedImage = double(originalImage) + 2 * unsharpMask;
```

% Convert the results back to uint8 for display

```
unsharpMaskedImage = uint8(unsharpMaskedImage);
```

```
highBoostMaskedImage = uint8(highBoostMaskedImage);  
  
% Display or save the sharpened images  
imshow(unsharpMaskedImage);  
imwrite(unsharpMaskedImage, 'unsharp_masked_image.jpg');  
imshow(highBoostMaskedImage);  
imwrite(highBoostMaskedImage, 'high_boost_image.jpg');
```

Output:

Conclusion: In this practical, I learnt about spatial domain operations, including cross-correlation, convolution, smoothing, sharpening, non-linear filtering, noise removal, unsharp masking, and high-boost filtering, demonstrating a wide range of image processing techniques.

PRACTICAL - 5

Aim: Implement code to

- i. convert images from space domain to frequency domain and observe their spectrum
- ii. observe aliasing in down-sampled images and apply anti-aliasing filter to reduce effect of aliasing
- iii. apply frequency domain low-pass filters of different types and cut-off frequencies on images and observe their effects
- iv. apply frequency domain high-pass filters of different types and cut-off frequencies on images and observe their effects
- v. add periodic noise on images in frequency domain, apply notch filters to remove noise and restore original image

Program:

i.

% Read the original image

```
originalImage = imread('jesse.png');
```

% Convert image to grayscale if it's a color image

```
if size(originalImage, 3) == 3
```

```
    originalImage = rgb2gray(originalImage);
```

```
end
```


% Perform Fourier Transform to convert to frequency domain

```
frequencyDomainImage = fftshift(fft2(originalImage));
```

% Calculate magnitude spectrum for visualization

```
magnitudeSpectrum = abs(frequencyDomainImage);
```

```
% Display the magnitude spectrum  
imshow(log(1 + magnitudeSpectrum), []);
```

Output:**Program:**

ii.

```
% Downsample the original image (introducing aliasing)
```

```
downsampledImage = imresize(originalImage, 0.5);
```

```
% Apply anti-aliasing filter (Gaussian filter in this case)
```


```
filteredImage = imgaussfilt(downsampledImage, 1);
```

```
% Display the down-sampled and filtered images
```

```
figure;
```

```
subplot(1, 2, 1), imshow(downsampledImage), title('Down-sampled Image');
```

```
subplot(1, 2, 2), imshow(filteredImage), title('Filtered Image');
```

Output:**Program:**

iii.

% Apply frequency domain low-pass filters (e.g., Gaussian filter)

cutoffFrequency = 50;

lowpassFilteredImage = frequencyDomainImage;


lowpassFilteredImage(abs(lowpassFilteredImage) > cutoffFrequency) = 0;

% Inverse Fourier Transform to obtain filtered image

filteredImage = abs(iff2(fftshift(lowpassFilteredImage))));

% Display the filtered image

imshow(filteredImage, []);

Output:

Program:

iv.

% Apply frequency domain high-pass filters (e.g., Ideal high-pass filter)

cutoffFrequency = 50;

highpassFilteredImage = frequencyDomainImage;

highpassFilteredImage(abs(highpassFilteredImage) < cutoffFrequency) = 0;

% Inverse Fourier Transform to obtain filtered image

filteredImage = abs(fft2(fftshift(highpassFilteredImage))));

% Display the filtered image

imshow(filteredImage, []);

Output:

Program:

v.

% Add periodic noise in frequency domain

noiseAmplitude = 20;

noiseFrequencyX = 50;

noiseFrequencyY = 30;

noisyFrequencyDomainImage = frequencyDomainImage;

noisyFrequencyDomainImage(noiseFrequencyY, noiseFrequencyX) = ...

noisyFrequencyDomainImage(noiseFrequencyY, noiseFrequencyX) +
noiseAmplitude;

% Apply notch filters to remove noise

notchFilterRadius = 10;

notchFilter = ones(size(noisyFrequencyDomainImage));

notchFilter(noiseFrequencyY-
notchFilterRadius:noiseFrequencyY+notchFilterRadius, ...

noiseFrequencyX-notchFilterRadius:noiseFrequencyX+notchFilterRadius) = 0;
denoisedFrequencyDomainImage = noisyFrequencyDomainImage .* notchFilter;

% Inverse Fourier Transform to obtain denoised image


denoisedImage = abs(iff2(iff2shift(denoisedFrequencyDomainImage)));

% Display the noisy and denoised images

figure;

subplot(1, 2, 1), imshow(abs(iff2(iff2shift(noisyFrequencyDomainImage))), []),
title('Noisy Image');

subplot(1, 2, 2), imshow(denoisedImage, []), title('Denoised Image');

Output:

Conclusion: In this practical, I learnt about aliasing and removing noise from images.

PRACTICAL - 6

Aim: Implement code to detect

- i. edges in different images using laplacian operator
- ii. edges in different images using sobel operator
- iii. edges in different images using prewitt operator
- iv. edges in different images using canny operator
- v. lines in different images using Hough Transform

Program:

i.

% Apply Laplacian operator for edge detection

```
laplacianEdges = edge(grayImage, 'log');
```

% Display or save the edges detected using Laplacian operator

```
imshow(laplacianEdges);
```

```
imwrite(laplacianEdges, 'laplacian_edges.jpg');
```

Output:



Program:

ii.


```
% Apply Sobel operator for edge detection (horizontal and vertical edges)
```

```
sobelEdges = edge(grayImage, 'sobel');
```

```
% Display or save the edges detected using Sobel operator
```

```
imshow(sobelEdges);
```

```
imwrite(sobelEdges, 'sobel_edges.jpg');
```

Output:**Program:**

iii.


```
% Apply Prewitt operator for edge detection (horizontal and vertical edges)
```

```
prewittEdges = edge(grayImage, 'prewitt');
```

```
% Display or save the edges detected using Prewitt operator
```

```
imshow(prewittEdges);
```

```
imwrite(prewittEdges, 'prewitt_edges.jpg');
```

Output:**Program:**

iv.

% Apply Canny operator for edge detection

```
cannyEdges = edge(grayImage, 'canny');
```

% Display or save the edges detected using Canny operator

```
imshow(cannyEdges);
```

```
imwrite(cannyEdges, 'canny_edges.jpg');
```

Output:

Program:

v.

% Apply edge detection (e.g., Canny) for better Hough transform results

edges = edge(grayImage, 'canny');

% Perform Hough Transform for line detection

[H,theta,rho] = hough(edges);

peaks = houghpeaks(H, 10); % Specify the number of peaks to detect

% Find lines in the image

lines = houghlines(edges, theta, rho, peaks);

% Draw lines on the original color image

figure, imshow(originalImage), hold on

for k = 1:length(lines)

 xy = [lines(k).point1; lines(k).point2];

 plot(xy(:,1), xy(:,2), 'LineWidth', 2, 'Color', 'r');

end

hold off;

% Display or save the image with detected lines

imwrite(originalImage, 'lines_detected_image.jpg');

Output:

Conclusion: In this practical, I implemented various edge detection operators, including Laplacian, Sobel, Prewitt, and Canny, as well as the Hough Transform for line detection.

PRACTICAL - 7

Aim: Implement code to detect features using

- i. Harris corner detector
- ii. Shi-Tomasi corner detector
- iii. Scale Invariant Feature Transform (SIFT)
- iv. Speeded up Robust Feature (SURF)
- v. Oriented FAST and Rotated BRIEF (ORB)

Program:

```
import cv2

import numpy as np

import matplotlib.pyplot as plt

image_path = "/content/pexels-brett-sayles-6424244 (1).jpg"
image = cv2.imread(image_path)
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Detect corners using Harris corner detector
corner_image = cv2.cornerHarris(gray_image, blockSize=2, ksize=3, k=0.04)

# Threshold the corner response to identify strong corners
threshold = 0.01 * corner_image.max()
corner_image[corner_image < threshold] = 0

# Dilate the corners to make them more visible
corner_image_dilated = cv2.dilate(corner_image, None)
```

```
# Mark detected corners on the original image
```

```
image_with_corners = image.copy()
```

```
image_with_corners[corner_image_dilated > 0.01 * corner_image_dilated.max()]  
= [0, 0, 255] # Red color
```

```
# Display the original image with detected corners
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

```
plt.subplot(122), plt.imshow(cv2.cvtColor(image_with_corners,  
cv2.COLOR_BGR2RGB)), plt.title('Image with Detected Corners')
```

```
plt.show()
```

```
# Detect corners using Shi-Tomasi corner detector
```

```
corners = cv2.goodFeaturesToTrack(gray_image, maxCorners=100,  
qualityLevel=0.01, minDistance=10)
```

```
# Convert corners to integer coordinates
```

```
corners = np.int0(corners)
```

```
# Draw detected corners on the original image
```

```
image_with_corners = image.copy()
```

```
for corner in corners:
```

```
    x, y = corner.ravel()
```

```
    cv2.circle(image_with_corners, (x, y), 3, 255, -1) # Draw a circle at each corner
```

```
# Display the original image with detected corners
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

```
plt.subplot(122), plt.imshow(cv2.cvtColor(image_with_corners,  
cv2.COLOR_BGR2RGB)), plt.title('Image with Detected Corners')
```

```
plt.show()
```

```
# Create an SIFT object
```

```
sift = cv2.SIFT_create()
```

```
# Detect keypoints and compute descriptors
```

```
keypoints, descriptors = sift.detectAndCompute(gray_image, None)
```

```
# Draw detected keypoints on the original image
```

```
image_with_keypoints = cv2.drawKeypoints(image, keypoints, None)
```

```
# Display the original image with detected keypoints
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

```
plt.subplot(122), plt.imshow(cv2.cvtColor(image_with_keypoints,  
cv2.COLOR_BGR2RGB)), plt.title('Image with Detected Keypoints')
```

```
plt.show()
```

```
# Create a SURF object
```

```
surf = cv2.SURF_create()
```



```
# Detect keypoints and compute descriptors
```

```
keypoints, descriptors = surf.detectAndCompute(gray_image, None)
```

```
# Draw detected keypoints on the original image
```

```
image_with_keypoints = cv2.drawKeypoints(image, keypoints, None, (0, 255, 0),  
4)
```

```
# Display the original image with detected keypoints
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

```
plt.subplot(122), plt.imshow(cv2.cvtColor(image_with_keypoints,  
cv2.COLOR_BGR2RGB)), plt.title('Image with Detected Keypoints')
```

```
plt.show()
```

```
# Create an ORB object
```

```
orb = cv2.ORB_create()
```

```
# Detect keypoints and compute descriptors
```

```
keypoints, descriptors = orb.detectAndCompute(gray_image, None)
```

```
# Draw detected keypoints on the original image
```

```
image_with_keypoints = cv2.drawKeypoints(image, keypoints, None, (0, 255, 0),  
4)
```

```
# Display the original image with detected keypoints
```


```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

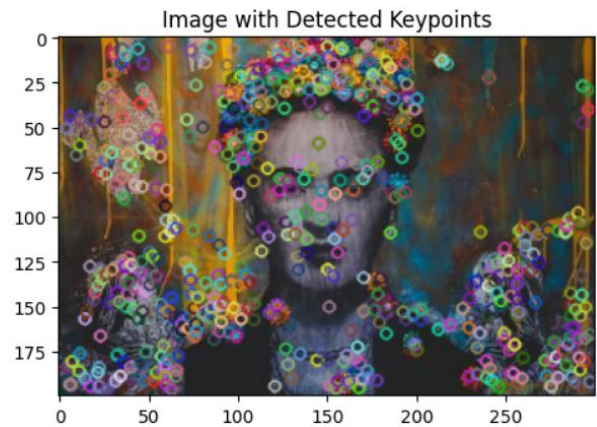
```
plt.subplot(122), plt.imshow(cv2.cvtColor(image_with_keypoints,  
cv2.COLOR_BGR2RGB)), plt.title('Image with Detected Keypoints')
```

```
plt.show()
```

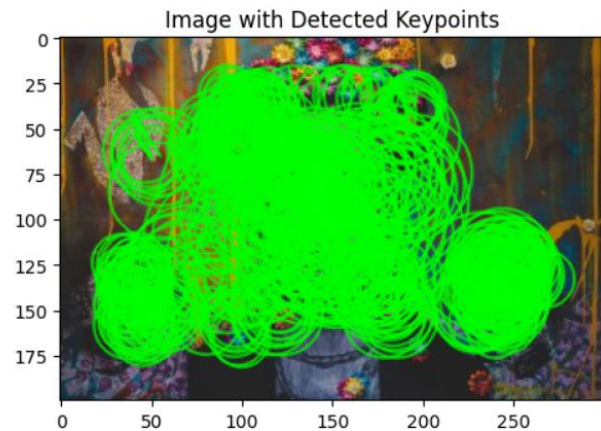
Output:



iii



v



Conclusion: In this practical, I learnt various feature detection methods, including Harris and Shi-Tomasi corner detectors, SIFT, SURF, and ORB, for feature-based image analysis.

PRACTICAL - 8

Aim: Implement a code to segment an image

- i. of mutually touching coins using distance transform along with watershed algorithms.
- ii. using the K-means algorithm.
- iii. using the Grabcut algorithm.

Program:

```
import cv2

import numpy as np

import matplotlib.pyplot as plt

image_path = "/content/pexels-brett-sayles-6424244 (1).jpg"
image = cv2.imread(image_path)
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Threshold the grayscale image to create a binary mask of the coins
_, binary_mask = cv2.threshold(gray_image, 0, 255, cv2.THRESH_BINARY_INV
+ cv2.THRESH_OTSU)

# Perform morphological operations to remove noise and separate touching coins
kernel = np.ones((3, 3), np.uint8)

opening = cv2.morphologyEx(binary_mask, cv2.MORPH_OPEN, kernel,
iterations=2)

# Calculate the distance transform
dist_transform = cv2.distanceTransform(opening, cv2.DIST_L2, 5)
```

```
_, sure_foreground = cv2.threshold(dist_transform, 0.7 * dist_transform.max(),  
255, 0)
```

```
# Find sure background
```

```
sure_background = cv2.dilate(opening, kernel, iterations=3)
```

```
# Subtract sure background from sure foreground to get unknown region
```

```
sure_foreground = np.uint8(sure_foreground)
```

```
unknown = cv2.subtract(sure_background, sure_foreground)
```

```
# Label markers for watershed
```

```
_, markers = cv2.connectedComponents(sure_foreground)
```

```
# Add 1 to all labels to ensure that sure background is not 0 (unlabeled)
```

```
markers = markers + 1
```

```
markers[unknown == 255] = 0
```

```
# Apply the watershed algorithm
```

```
cv2.watershed(image, markers)
```

```
image[markers == -1] = [0, 0, 255] # Mark segmented regions in red
```

```
# Display the original image with segmented coins
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(121), plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)),  
plt.title('Original Image')
```

```
plt.subplot(122), plt.imshow(markers, cmap='tab20'), plt.title('Segmented  
Regions')
```

```
plt.show()
```

```
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert to RGB
```

```
# Reshape the image to a 2D array of pixels
```

```
pixels = image.reshape(-1, 3)
```

```
# Define the number of clusters (K)
```

```
K = 3
```

```
# Apply K-means clustering
```

```
criteria = (cv2.TERM_CRITERIA_EPS + cv2.TERM_CRITERIA_MAX_ITER,  
100, 0.2)
```

```
_, labels, centers = cv2.kmeans(np.float32(pixels), K, None, criteria, 10,  
cv2.KMEANS_RANDOM_CENTERS)
```

```
# Convert the labels to 8-bit for visualization
```

```
labels = labels.reshape(image.shape[0], image.shape[1]).astype(np.uint8)
```

```
# Create a mask for each segment
```

```
segmented_images = []
```

```
for i in range(K):
```

```
    mask = np.where(labels == i, 255, 0).astype(np.uint8)
```

```
    segmented_images.append(mask)
```

```
# Display the original image and segmented images
plt.figure(figsize=(12, 6))
plt.subplot(131), plt.imshow(image), plt.title('Original Image')
plt.subplot(132), plt.imshow(segmented_images[0], cmap='gray'),
plt.title('Segment 1')
plt.subplot(133), plt.imshow(segmented_images[1], cmap='gray'),
plt.title('Segment 2')
plt.show()

# Create a mask and initialize it with zeros
mask = np.zeros(image.shape[:2], np.uint8)

# Define a rectangular region of interest (ROI) for initialization
rect = (50, 50, image.shape[1] - 50, image.shape[0] - 50)

# Initialize the GrabCut algorithm with the image, mask, and ROI
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)
cv2.grabCut(image, mask, rect, bgdModel, fgdModel, 5,
cv2.GC_INIT_WITH_RECT)

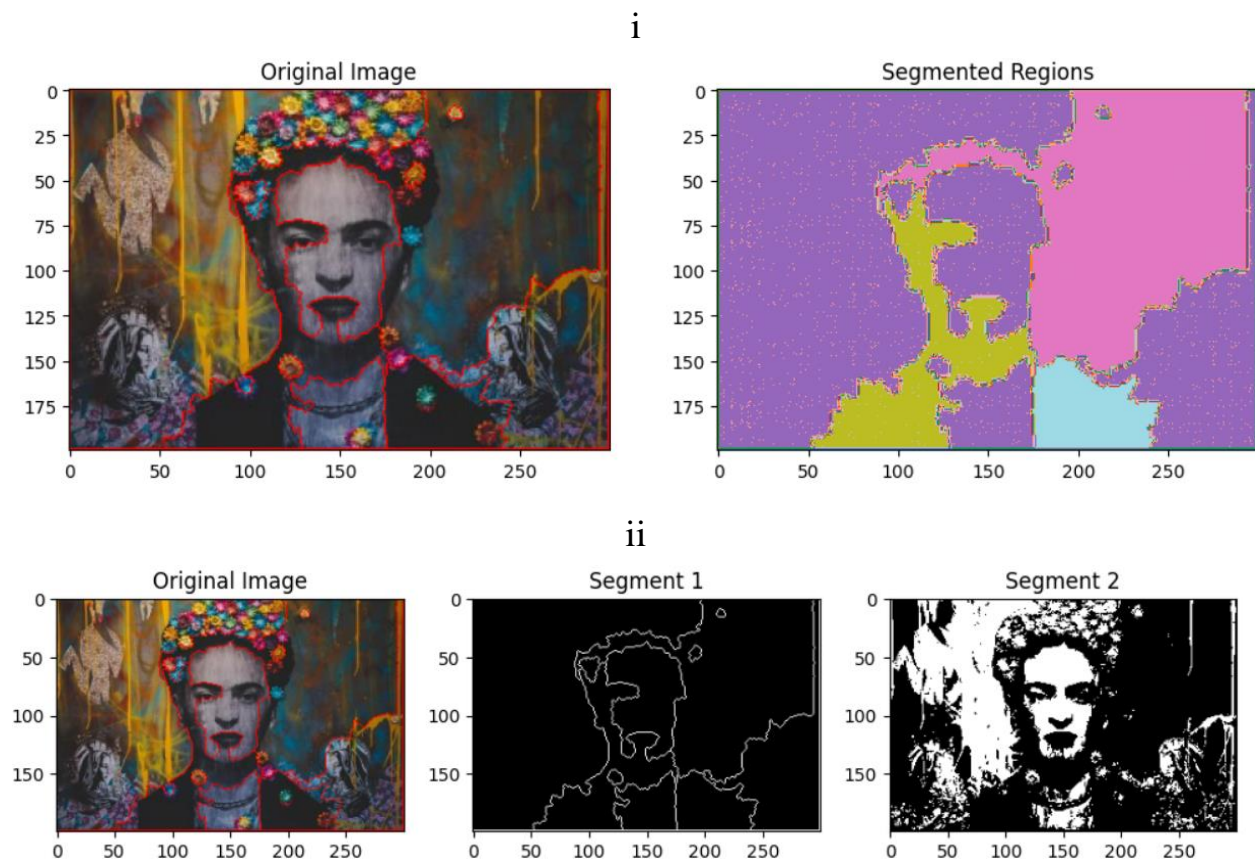
# Modify the mask to create a binary mask for the foreground
mask2 = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')

# Multiply the original image with the binary mask to extract the segmented object
```




```
segmented_image = image * mask2[:, :, np.newaxis]

# Display the original image and the segmented result
plt.figure(figsize=(12, 6))
plt.subplot(121), plt.imshow(image), plt.title('Original Image')
plt.subplot(122), plt.imshow(segmented_image), plt.title('Segmented Image')
plt.show()
```

Output:

iii



Conclusion: In this practical, I segmented images of mutually touching coins by employing distance transform with watershed, K-means, and Grabcut algorithms

PRACTICAL - 9

Aim: i. Implement face detection and eye detection using HAAR cascade classifiers.
ii. Implement face detection using Viola Jones method and Adaboost training algorithm.
iii. Implement car detection and pedestrian detection using HAAR cascade classifiers.

Program:

```
import cv2

from google.colab.patches import cv2_imshow # Import cv2_imshow for Colab
import matplotlib.pyplot as plt

# Load pre-trained HAAR cascade classifiers for face and eye detection
face_cascade = cv2.CascadeClassifier(cv2.data.harcascades +
'haarcascade_frontalface_default.xml')

eye_cascade = cv2.CascadeClassifier(cv2.data.harcascades +
'haarcascade_eye.xml')

# Load an image
image_path = "/content/saulgoodman.png"
image = cv2.imread(image_path)
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Detect faces in the image
faces = face_cascade.detectMultiScale(gray_image, scaleFactor=1.3,
minNeighbors=5, minSize=(30, 30))
```

```
# Iterate over detected faces and draw rectangles around them
```

```
for (x, y, w, h) in faces:
```

```
    cv2.rectangle(image, (x, y), (x + w, y + h), (255, 0, 0), 2)
```

```
    roi_gray = gray_image[y:y + h, x:x + w]
```

```
    roi_color = image[y:y + h, x:x + w]
```

```
# Detect eyes within each face region
```

```
eyes = eye_cascade.detectMultiScale(roi_gray)
```

```
for (ex, ey, ew, eh) in eyes:
```

```
    cv2.rectangle(roi_color, (ex, ey), (ex + ew, ey + eh), (0, 255, 0), 2)
```

```
# Display the image with detected faces and eyes using cv2_imshow
```

```
cv2_imshow(image)
```

```
# Detect faces in the image
```

```
faces = face_cascade.detectMultiScale(image, scaleFactor=1.3, minNeighbors=5,  
minSize=(30, 30))
```

```
# Draw rectangles around detected faces
```

```
for (x, y, w, h) in faces:
```

```
    cv2.rectangle(image, (x, y), (x + w, y + h), (255, 0, 0), 2)
```

```
# Display the image with detected faces
```

```
cv2_imshow(image)
```

```
cv2.waitKey(0)
```

```
cv2.destroyAllWindows()
```

```
# Specify the paths to the Haar cascade XML files
```

```
car_cascade_path = '/content/haarcascade_car.xml'
```

```
pedestrian_cascade_path = '/content/haarcascade_fullbody.xml'
```

```
# Load the Haar cascade classifiers for car and pedestrian detection
```

```
car_cascade = cv2.CascadeClassifier(car_cascade_path)
```

```
pedestrian_cascade = cv2.CascadeClassifier(pedestrian_cascade_path)
```

```
# Load an image for detection
```

```
image_path = "/content/anthony-rosset-YLaLy6wIDiY-unsplash.jpg"
```

```
image = cv2.imread(image_path)
```

```
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
```

```
# Detect cars in the image
```

```
cars = car_cascade.detectMultiScale(gray_image, scaleFactor=1.1,  
minNeighbors=5, minSize=(20, 20))
```

```
# Detect pedestrians in the image
```

```
pedestrians = pedestrian_cascade.detectMultiScale(gray_image, scaleFactor=1.1,  
minNeighbors=5, minSize=(20, 20))
```

```
# Brighter colors for drawing rectangles
```

```
car_color = (0, 0, 255) # Bright red (BGR color format)
```

```
pedestrian_color = (0, 255, 0) # Bright green (BGR color format)
```

Line thickness for drawing rectangles

line_thickness = 8 # Adjust as needed for thicker borders

Loop through the detected cars and draw rectangles with thicker borders

for (x, y, w, h) in cars:

cv2.rectangle(image, (x, y), (x + w, y + h), car_color, line_thickness)

Loop through the detected pedestrians and draw rectangles with thicker borders

for (x, y, w, h) in pedestrians:

cv2.rectangle(image, (x, y), (x + w, y + h), pedestrian_color, line_thickness)

Convert the image from BGR to RGB for displaying with matplotlib

image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

Display the image with detected cars and pedestrians


plt.imshow(image_rgb)

plt.axis('off') # Turn off axis labels

plt.show()

Output:


i



ii



iii





Conclusion: In this practical, I applied HAAR cascade classifiers for face, eye, car, and pedestrian detection, demonstrating robust object detection techniques in different contexts.

PRACTICAL - 10

- Aim:** i. Implement code to perform feature extraction on given images of faces using Histogram of Gradients.
ii. Implement code to apply Principal Component Analysis on extracted features in objective i.
iii. Implement code to recognize faces using the SVM classifier.

Program:

```
import matplotlib.pyplot as plt
from skimage import io, feature
from skimage.color import rgb2gray
from sklearn.decomposition import PCA
import numpy as np

# Load an image of a face (replace with your image path)
image_path = "/content/download (1).jpeg"
image = io.imread(image_path)

# Convert the image to grayscale
gray_image = rgb2gray(image)

# Compute HOG features
hog_features, hog_image = feature.hog(gray_image, visualize=True)

# Display the original image
```

```
plt.figure(figsize=(8, 4))
plt.subplot(121)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')

# Display the HOG image
plt.subplot(122)
plt.imshow(hog_image, cmap='gray')
plt.title('HOG Features')
plt.axis('off')

plt.tight_layout()
plt.show()

# Print the HOG feature vector (histogram)
print("HOG Feature Vector (Histogram):")
print(hog_features)

# Extracted HOG features (replace with your own features)
hog_features = np.array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])

# Create a PCA instance without specifying the number of components
pca = PCA()
```

```
# Fit the PCA model to the HOG features
pca.fit(hog_features.reshape(-1, 1)) # Reshape to a single feature per sample

# Transform the features into the PCA space
hog_features_pca = pca.transform(hog_features.reshape(-1, 1))

# Explained variance ratio of all components
explained_variance_ratio = pca.explained_variance_ratio_

# Print the transformed features and explained variance
print("Transformed HOG Features (PCA):")
print(hog_features_pca)
print("Explained Variance Ratio:")
print(explained_variance_ratio)
```

iii)

```
from sklearn.datasets import fetch_lfw_people
faces = fetch_lfw_people(min_faces_per_person=60)
```

```
faces.DESCR
```

```
import matplotlib.pyplot as plt
```

```
fig, splts = plt.subplots(2, 4)
for i, splts in enumerate(splts.flat):
```

```
splits.imshow(faces.images[i], cmap='magma')
splits.set(xticks=[], yticks=[],
           xlabel=faces.target_names[faces.target[i]])

from sklearn.model_selection import train_test_split

X = faces.data
y = faces.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
                                                    random_state=42)

from sklearn.svm import SVC
from sklearn.decomposition import PCA as RandomizedPCA
from sklearn.pipeline import make_pipeline

# For dimensionality reduction
pca = RandomizedPCA(n_components=150, whiten=True, random_state=42)
svc = SVC(kernel='rbf', class_weight='balanced')
model = make_pipeline(pca, svc)

model.fit(X_train, y_train)

from sklearn.metrics import accuracy_score

predictions = model.predict(X_test)
```

```
accuracy_score(predictions, y_test)
```

```
from colorama import Fore
```

```
incorrect = 0
```

```
length = len(predictions)
```

```
print("Actual\t\t\tPredicted\n")
```

```
for i in range(len(predictions)):
```

```
    if predictions[i] != y_test[i]: # if predictions and actual values are not equal
```

```
        prediction_name = faces.target_names[predictions[predictions[i]]] # Getting  
the predicted name
```

```
        actual_name = faces.target_names[y_test[y_test[i]]] # Getting the actual name
```

```
        incorrect+=1
```

```
        print("{}\t\t{}".format(Fore.GREEN + actual_name,  
Fore.RED+prediction_name))
```

```
print("{} {} are classified as correct and {} {} are classified as incorrect!".format(length-  
incorrect, incorrect))
```


Output:

i

Original Image




HOG Features



HOG Feature Vector (Histogram):

```
[0.      0.      0.      ... 0.01053878 0.00620127 0.24241371]
```

ii

 Transformed HOG Features (PCA):

```
[[ 0.45]
 [ 0.35]
 [ 0.25]
 [ 0.15]
 [ 0.05]
 [-0.05]
 [-0.15]
 [-0.25]
 [-0.35]
 [-0.45]]
```

Explained Variance Ratio:

```
[1.]
```

iii



Colin Powell



George W Bush



George W Bush



George W Bush



Hugo Chavez



George W Bush



Junichiro Koizumi



George W Bush

Actual	Predicted
Junichiro Koizumi	Junichiro Koizumi
George W Bush	George W Bush
Junichiro Koizumi	George W Bush
Colin Powell	Junichiro Koizumi
George W Bush	Junichiro Koizumi
Colin Powell	Junichiro Koizumi
Colin Powell	George W Bush
George W Bush	Junichiro Koizumi
George W Bush	George W Bush
George W Bush	Junichiro Koizumi
Junichiro Koizumi	George W Bush
Colin Powell	Junichiro Koizumi
Junichiro Koizumi	George W Bush
Colin Powell	George W Bush
Junichiro Koizumi	George W Bush
George W Bush	George W Bush
George W Bush	Junichiro Koizumi
Junichiro Koizumi	George W Bush
George W Bush	George W Bush
Junichiro Koizumi	George W Bush
Junichiro Koizumi	George W Bush
George W Bush	George W Bush
George W Bush	Junichiro Koizumi
George W Bush	Junichiro Koizumi
George W Bush	Junichiro Koizumi
George W Bush	Junichiro Koizumi
George W Bush	Junichiro Koizumi
Colin Powell	Junichiro Koizumi
Junichiro Koizumi	George W Bush
Colin Powell	George W Bush
Junichiro Koizumi	Junichiro Koizumi
Gerhard Schroeder	Junichiro Koizumi
George W Bush	Junichiro Koizumi
Colin Powell	Junichiro Koizumi

436 are classified as correct and 104 are classified as incorrect!

Conclusion: In this practical, I performed facial feature extraction with Histogram of Gradients, applied Principal Component Analysis (PCA) on the features, and achieved facial recognition using SVM classifiers.

PRACTICAL - 11

Aim:

- i. Implement code to extract facial landmarks on given images.
- ii. Implement code to merge faces (face swaps) using extracted facial landmark features on given images.
- iii. Implement code to merge faces (face swaps) using extracted facial landmark features on live video.

Program:

```
import cv2
import dlib
import numpy as np
from google.colab.patches import cv2_imshow

# Load the face detector from dlib (HOG-based)
face_detector = dlib.get_frontal_face_detector()

# Load the facial landmarks predictor from dlib
landmark_predictor =
dlib.shape_predictor("/content/drive/MyDrive/shape_predictor_68_face_landmarks
.dat")

# Load an image
image_path = "/content/saulgoodman.png"
image = cv2.imread(image_path)

# Convert the image to grayscale (required by dlib)
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
```

```
# Detect faces in the image
faces = face_detector(gray_image)

# Loop over the detected faces
for face in faces:
    # Detect facial landmarks
    landmarks = landmark_predictor(gray_image, face)

    # Loop over the facial landmarks and draw them on the image
    for i in range(68):
        x, y = landmarks.part(i).x, landmarks.part(i).y
        cv2.circle(image, (x, y), 2, (0, 255, 0), -1) # Draw a green circle at each
        landmark point

# Display the image with facial landmarks
cv2_imshow(image)

# Load the source and target images
source_image = cv2.imread("/content/jesse.png")
target_image = cv2.imread("/content/download (1).jpeg")

# Detect faces in both images
source_faces = face_detector(source_image)
target_faces = face_detector(target_image)

# Ensure one face is detected in each image
if len(source_faces) != 1 or len(target_faces) != 1:
    print("Error: Exactly one face must be present in each image.")
else:
```

```
# Get facial landmarks for both faces
source_landmarks = landmark_predictor(source_image, source_faces[0])
target_landmarks = landmark_predictor(target_image, target_faces[0])

# Convert landmarks to NumPy arrays
source_landmarks = np.array([[p.x, p.y] for p in source_landmarks.parts()])
target_landmarks = np.array([[p.x, p.y] for p in target_landmarks.parts()])

# Compute the affine transformation matrix
transformation_matrix, _ = cv2.estimateAffinePartial2D(source_landmarks,
target_landmarks)


# Warp the source face to match the target face
warped_face = cv2.warpAffine(source_image, transformation_matrix,
(target_image.shape[1], target_image.shape[0]))

# Blend the warped face onto the target face
alpha = 0.7 # Adjust this value for blending intensity
beta = 1.0 - alpha
blended_face = cv2.addWeighted(target_image, alpha, warped_face, beta, 0)

# Save or display the resulting image
cv2.imwrite("output_face_swap.jpg", blended_face)
cv2.imshow( blended_face)
```

Output:

i



ii



Conclusion: In this practical, I extracted facial landmarks, performed face swaps using landmark features on static images, and extended the capability to live video.

PRACTICAL - 12

Aim: Implement Deep Learning concepts (using DIGITS/TensorFlow/Pytorch)

- i. Image Classification
- ii. Image Segmentation
- iii. Object Detection
- iv. Transfer Learning
- v. Face Recognition
- vi. Emotion Recognition

Program:

Output:

i

```
from google.colab import drive
drive.mount('/content/drive')
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
# Define your dataset directory on Google Colab
dataset_dir = '/content/drive/MyDrive/practical12/traindata/'
# Define hyperparameters
batch_size = 32
epochs = 10
input_shape = (224, 224, 3) # Adjust the input shape according to your images
num_classes = 2 # Change this to the number of classes in your dataset
# Data augmentation and preprocessing
train_datagen = ImageDataGenerator(
```

```
    rescale=1.0/255.0,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True,
    shear_range=0.2,
    zoom_range=0.2
)
# Create a generator for training data
train_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=input_shape[:2],
    batch_size=batch_size,
    class_mode='categorical', # Use 'binary' for binary classification
    shuffle=True
)
# Build a convolutional neural network (CNN) model
model = keras.Sequential([
    keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=input_shape),
    keras.layers.MaxPooling2D((2, 2)),
    keras.layers.Conv2D(64, (3, 3), activation='relu'),
    keras.layers.MaxPooling2D((2, 2)),
    keras.layers.Conv2D(128, (3, 3), activation='relu'),
    keras.layers.MaxPooling2D((2, 2)),
    keras.layers.Flatten(),
```

```
keras.layers.Dense(128, activation='relu'),
keras.layers.Dense(num_classes, activation='softmax')
])
# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
# Train the model
history = model.fit(
    train_generator,
    steps_per_epoch=len(train_generator),
    epochs=epochs
)
# Save the trained model
model.save('image_classification_model.h5')
# Optionally, save training history for analysis
import pickle
with open('training_history.pkl', 'wb') as file:
    pickle.dump(history.history, file)
import os
import numpy as np
from tensorflow.keras.preprocessing import image
# Load the trained model
model = keras.models.load_model('image_classification_model.h5')
# Define a function to predict an individual image
```

```
def predict_image(image_path):
    img = image.load_img(image_path, target_size=(224, 224))
    img = image.img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = img / 255.0
    predictions = model.predict(img)
    class_index = np.argmax(predictions)
    return class_index

# Define the directory containing test images on Google Colab
test_dir = '/content/drive/MyDrive/practical12/traindata/testimage/'

# Loop through test images and make predictions
for filename in os.listdir(test_dir):
    if filename.endswith('.jpg'):
        image_path = os.path.join(test_dir, filename)
        class_index = predict_image(image_path)
        print(f"Image: {filename}, Predicted Class Index: {class_index}")

from google.colab import drive
drive.mount('/content/drive')
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Define your dataset directory on Google Colab
dataset_dir = '/content/drive/MyDrive/practical12/traindata/'

# Define hyperparameters
```



```
batch_size = 32
epochs = 10
input_shape = (224, 224, 3) # Adjust the input shape according to your images
num_classes = 2 # Change this to the number of classes in your dataset
# Data augmentation and preprocessing
train_datagen = ImageDataGenerator(
    rescale=1.0/255.0,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True,
    shear_range=0.2,
    zoom_range=0.2
)
# Create a generator for training data
train_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=input_shape[:2],
    batch_size=batch_size,
    class_mode='categorical', # Use 'binary' for binary classification
    shuffle=True
)
# Build a convolutional neural network (CNN) model
model = keras.Sequential([
    keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=input_shape),
```

```
keras.layers.MaxPooling2D((2, 2)),
keras.layers.Conv2D(64, (3, 3), activation='relu'),
keras.layers.MaxPooling2D((2, 2)),
keras.layers.Conv2D(128, (3, 3), activation='relu'),
keras.layers.MaxPooling2D((2, 2)),
keras.layers.Flatten(),
keras.layers.Dense(128, activation='relu'),
keras.layers.Dense(num_classes, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(
    train_generator,
    steps_per_epoch=len(train_generator),
    epochs=epochs
)

# Save the trained model
model.save('image_classification_model.h5')

# Optionally, save training history for analysis
import pickle
with open('training_history.pkl', 'wb') as file:
    pickle.dump(history.history, file)
```

```
import os
import numpy as np
from tensorflow.keras.preprocessing import image
# Load the trained model
model = keras.models.load_model('image_classification_model.h5')
# Define a function to predict an individual image
def predict_image(image_path):
    img = image.load_img(image_path, target_size=(224, 224))
    img = image.img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = img / 255.0
    predictions = model.predict(img)
    class_index = np.argmax(predictions)
    return class_index

# Define the directory containing test images on Google Colab
test_dir = '/content/drive/MyDrive/practical12/traindata/testimage/'

# Loop through test images and make predictions
for filename in os.listdir(test_dir):
    if filename.endswith('.jpg'):
        image_path = os.path.join(test_dir, filename)
        class_index = predict_image(image_path)
        print(f"Image: {filename}, Predicted Class Index: {class_index}")
```

ii.

Setting the dataset path

```
import pathlib
```

```
data_dir = pathlib.Path('/content/gdrive/MyDrive/Segmentation/dataset1')
```

```
image_count = len(list(data_dir.glob('*/*.png')))
```

```
print(image_count)
```

```
dir_data = "/content/gdrive/MyDrive/Segmentation/dataset1/"
```

```
dir_seg = dir_data + "/annotations_prepped_train/"
```

```
dir_img = dir_data + "/images_prepped_train/"
```

```
import cv2, os
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
## seaborn has white grid by default so I will get rid of this.
```

```
sns.set_style("whitegrid", {'axes.grid' : False})
```

```
ldseg = np.array(os.listdir(dir_seg))
```

```
## pick the first image file
```

```
fnm = ldseg[0]
```

```
print(fnm)
```

```
## read in the original image and segmentation labels
```

```
seg = cv2.imread(dir_seg + fnm ) # (360, 480, 3)
```

```
img_is = cv2.imread(dir_img + fnm )
```

```
print("seg.shape={ }, img_is.shape={ }".format(seg.shape,img_is.shape))
```

```
## Check the number of labels
mi, ma = np.min(seg), np.max(seg)
n_classes = ma - mi + 1

print("minimum seg = { }, maximum seg = { }, Total number of segmentation
classes = { }".format(mi,ma, n_classes))

fig = plt.figure(figsize=(5,5))
ax = fig.add_subplot(1,1,1)
ax.imshow(img_is)
ax.set_title("original image")
plt.show()

fig = plt.figure(figsize=(15,10))
for k in range(mi,ma+1):
    ax = fig.add_subplot(3, int(n_classes/3)+1, k+1)
    ax.imshow((seg == k)*1.0)
    ax.set_title("label = { }".format(k))
plt.show()

import random

def give_color_to_seg_img(seg,n_classes):
    """
    seg : (input_width,input_height,3)
    """
    if len(seg.shape)==3:
        seg = seg[:, :, 0]
    seg_img = np.zeros( (seg.shape[0],seg.shape[1],3) ).astype('float')
    colors = sns.color_palette("hls", n_classes)
    for c in range(n_classes):
```

```
    segc = (seg == c)
    seg_img[:, :, 0] += (segc * (colors[c][0]))
    seg_img[:, :, 1] += (segc * (colors[c][1]))
    seg_img[:, :, 2] += (segc * (colors[c][2]))
    return(seg_img)

input_height, input_width = 224, 224
output_height, output_width = 224, 224
ldseg = np.array(os.listdir(dir_seg))
for fnm in ldseg[np.random.choice(len(ldseg), 3, replace=False)]:
    fnm = fnm.split(".")[0]
    seg = cv2.imread(dir_seg + fnm + ".png") # (360, 480, 3)
    img_is = cv2.imread(dir_img + fnm + ".png")
    seg_img = give_color_to_seg_img(seg, n_classes)
    fig = plt.figure(figsize=(20, 40))
    ax = fig.add_subplot(1, 4, 1)
    ax.imshow(seg_img)
    ax = fig.add_subplot(1, 4, 2)
    ax.imshow(img_is/255.0)
    ax.set_title("original image {}".format(img_is.shape[:2]))
    ax = fig.add_subplot(1, 4, 3)
    ax.imshow(cv2.resize(seg_img, (input_height, input_width)))
    ax = fig.add_subplot(1, 4, 4)
    ax.imshow(cv2.resize(img_is, (output_height, output_width))/255.0)
    ax.set_title("resized to {}".format((output_height, output_width)))
    plt.show()
```

iv.

```
import tensorflow as tf

from tensorflow.keras.datasets import cifar10

from tensorflow.keras.applications import ResNet50

from tensorflow.keras.layers import GlobalAveragePooling2D, Dense

from tensorflow.keras.models import Model


# Load the CIFAR-10 dataset
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize pixel values to between 0 and 1
X_train, X_test = X_train / 255.0, X_test / 255.0

# Load the pre-trained ResNet50 model (excluding the top classification layers)
base_model = ResNet50(weights='imagenet', include_top=False)

# Add custom classification layers on top of the pre-trained model
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x) # Adjust the number of classes

# Create the transfer learning model
model = Model(inputs=base_model.input, outputs=predictions)

# Freeze the layers of the pre-trained model (optional)
for layer in base_model.layers:
    layer.trainable = False

# Compile the model
model.compile(optimizer='adam',
```

```
    loss='sparse_categorical_crossentropy', # Adjust the loss function
    metrics=['accuracy'])

# Train the model

history = model.fit(X_train, y_train,
                    epochs=10, # Adjust the number of epochs
                    validation_data=(X_test, y_test))

# Evaluate the model

test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f'Test accuracy: {test_accuracy:.4f}')
```

v.

```
#import OpenCV module
import cv2
import os
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

#function to detect face
def detect_face (img):
    #convert the test image to gray image
    gray = cv2.cvtColor (img, cv2.COLOR_BGR2GRAY)
    #load OpenCV face detector
    face_cas = cv2.CascadeClassifier ('-File name.xml-')
    faces = face_cas.detectMultiScale (gray, scaleFactor=1.3, minNeighbors=4);
```



```
#if no faces are detected then return image
if (len (faces) == 0):
    return None, None
#extract the face
faces [0]=(x, y, w, h)
#return only the face part
return gray[y: y+w, x: x+h], faces [0]

#this function will read all persons' training images, detect face #from each image
#and will return two lists of exactly same size, one list
def prepare_training_data(data_folder_path):

#-----STEP-1-----

#get the directories (one directory for each subject) in data folder
dirs = os.listdir(data_folder_path)
faces = []
labels = []
for dir_name in dirs:
    #our subject directories start with letter 's' so
    #ignore any non-relevant directories if any
    if not dir_name.startswith("s"):
        continue;

#-----STEP-2-----

#extract label number of subject from dir_name
#format of dir name = slabel
```

```
#, so removing letter 's' from dir_name will give us label
label = int(dir_name.replace("s", ""))

#build path of directory containin images for current subject subject
#sample subject_dir_path = "training-data/s1"
subject_dir_path = data_folder_path + "/" + dir_name

#get the images names that are inside the given subject directory
subject_images_names = os.listdir(subject_dir_path)

#-----STEP-3-----

#go through each image name, read image,
#detect face and add face to list of faces
for image_name in subject_images_names:
    #ignore system files like .DS_Store
    if image_name.startswith("."):
        continue;

    #build image path
    #sample image path = training-data/s1/1.pgm
    image_path = subject_dir_path + "/" + image_name

    #read image
    image = cv2.imread(image_path)

    #display an image window to show the image
    cv2.imshow("Training on image...", image)
    cv2.waitKey(100)

    #detect face
    face, rect = detect_face(image)
```

```
#-----STEP-4-----  
#we will ignore faces that are not detected  
if face is not None:  
    #add face to list of faces  
    faces.append(face)  
    #add label for this face  
    labels.append(label)  
cv2.destroyAllWindows()  
cv2.waitKey(1)  
cv2.destroyAllWindows()  
return faces, labels  
  
#let's first prepare our training data  
#data will be in two lists of same size  
#one list will contain all the faces  
#and other list will contain respective labels for each face  
print("Preparing data...")  
faces, labels = prepare_training_data("training-data")  
print("Data prepared")  
#print total faces and labels  
print("Total faces: ", len(faces))  
print("Total labels: ", len(labels))  
#create our LBPH face recognizer  
face_recognizer = cv2.face.createLBPHFaceRecognizer()  
#train our face recognizer of our training faces  
face_recognizer.train(faces, np.array(labels))
```

```
#function to draw rectangle on image
#according to given (x, y) coordinates and
#given width and heigh
def draw_rectangle(img, rect):
    (x, y, w, h) = rect
    cv2.rectangle(img, (x, y), (x+w, y+h), (0, 255, 0), 2)
#function to draw text on give image starting from
#passed (x, y) coordinates.
def draw_text(img, text, x, y):
    cv2.putText(img, text, (x, y), cv2.FONT_HERSHEY_PLAIN, 1.5, (0, 255, 0), 2)
#this function recognizes the person in image passed
#and draws a rectangle around detected face with name of the subject
def predict(test_img):
    #make a copy of the image as we don't want to chang original image
    img = test_img.copy()
    #detect face from the image
    face, rect = detect_face(img)
    #predict the image using our face recognizer
    label= face_recognizer.predict(face)
    #get name of respective label returned by face recognizer
    label_text = subjects[label]
    #draw a rectangle around face detected
    draw_rectangle(img, rect)
    #draw name of predicted person
    draw_text(img, label_text, rect[0], rect[1]-5)
```

```
return img

#load test images
test_img1 = cv2.imread("test-data/test1.jpg")
test_img2 = cv2.imread("test-data/test2.jpg")

#perform a prediction
predicted_img1 = predict(test_img1)
predicted_img2 = predict(test_img2)
print("Prediction complete")

#create a figure of 2 plots (one for each test image)
f, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 5))

#display test image1 result
ax1.imshow(cv2.cvtColor(predicted_img1, cv2.COLOR_BGR2RGB))

#display test image2 result
ax2.imshow(cv2.cvtColor(predicted_img2, cv2.COLOR_BGR2RGB))

#display both images
cv2.imshow("Tom cruise test", predicted_img1)
cv2.imshow("Shahrukh Khan test", predicted_img2)
cv2.waitKey(0)
cv2.destroyAllWindows()
cv2.waitKey(1)
cv2.destroyAllWindows()
```

vi.

read image

img = cv2.imread('img1.jpg')

call imshow() using plt object

plt.imshow(img[:, :, :-1])

display that image

plt.show()

storing the result

result = DeepFace.analyze(img,

actions = ['emotion'])

print result

print(result)

import the required modules

import cv2

import matplotlib.pyplot as plt

from deepface import DeepFace

read image

img = cv2.imread('img.jpg')

call imshow() using plt object

plt.imshow(img[:, :, :-1])

display that image

plt.show()

storing the result

result = DeepFace.analyze(img, actions=['emotion'])

print(result)

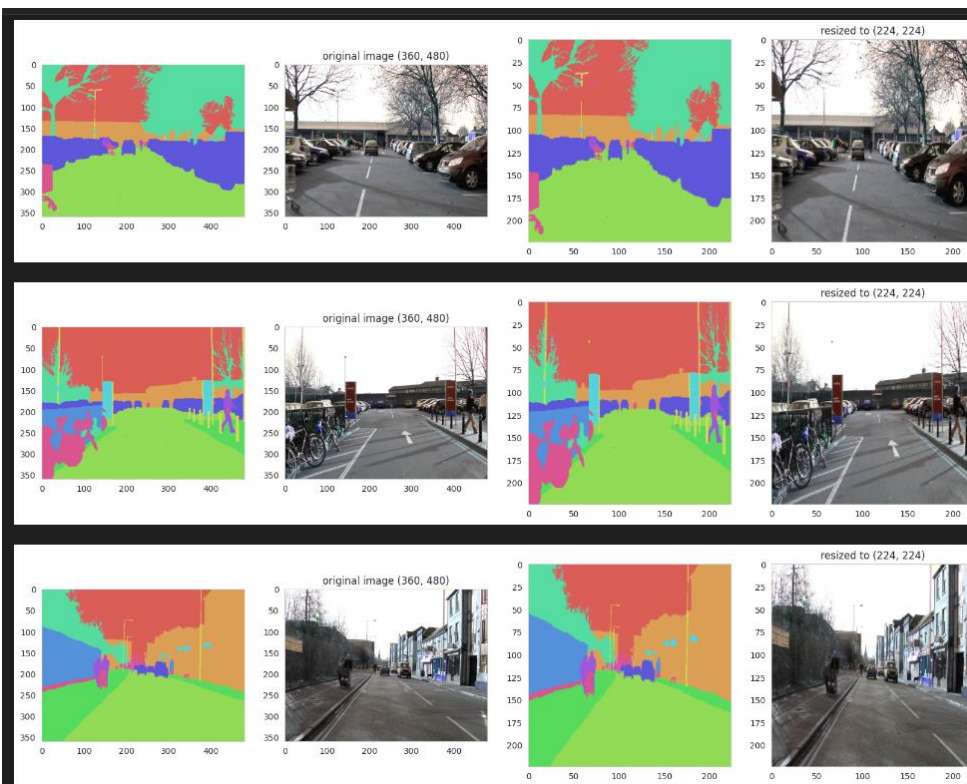
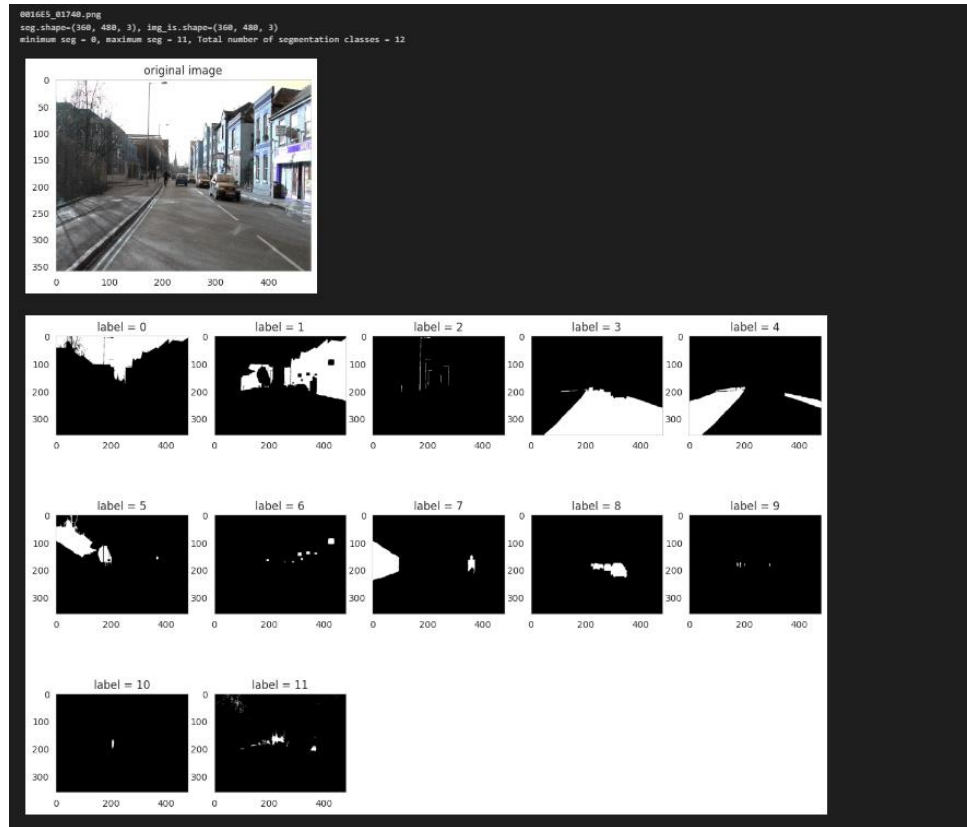
Output:

i.

```
Found 165 images belonging to 1 classes.
Epoch 1/10
6/6 [-----] - 51s 8s/step - loss: 1.8926 - accuracy: 0.1879
Epoch 2/10
6/6 [-----] - 3s 498ms/step - loss: 20.8844 - accuracy: 0.6121
Epoch 3/10
6/6 [-----] - 3s 394ms/step - loss: 132.8162 - accuracy: 0.5818
Epoch 4/10
6/6 [-----] - 3s 393ms/step - loss: 502.1111 - accuracy: 0.4182
Epoch 5/10
6/6 [-----] - 2s 381ms/step - loss: 765.7403 - accuracy: 0.5818
Epoch 6/10
6/6 [-----] - 3s 434ms/step - loss: 3444.7781 - accuracy: 0.4182
Epoch 7/10
6/6 [-----] - 4s 575ms/step - loss: 6166.1294 - accuracy: 0.4182
Epoch 8/10
6/6 [-----] - 3s 394ms/step - loss: 7368.7881 - accuracy: 0.6121
Epoch 9/10
6/6 [-----] - 3s 398ms/step - loss: 14062.2969 - accuracy: 0.4182
Epoch 10/10
6/6 [-----] - 3s 404ms/step - loss: 26948.4355 - accuracy: 0.7758
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3000: UserWarning: You are saving your model
  saving_api.save_model(
```

```
1/1 [-----] - 0s 70ms/step
Image: SPB_25-Mar_10-35-10_01.jpg, Predicted Class Index: 1
1/1 [-----] - 0s 20ms/step
Image: Robo_25-Mar_12-33-22_01.jpg, Predicted Class Index: 1
```

ii.



iv.

```

Epoch 1/10
1563/1563 [-----] - 37s 16ms/step - loss: 2.0382 - accuracy: 0.2593 - val_loss: 1.9041 - val_accuracy: 0.2964
Epoch 2/10
1563/1563 [-----] - 21s 13ms/step - loss: 1.8747 - accuracy: 0.3164 - val_loss: 1.8232 - val_accuracy: 0.3379
Epoch 3/10
1563/1563 [-----] - 20s 13ms/step - loss: 1.8194 - accuracy: 0.3401 - val_loss: 1.7811 - val_accuracy: 0.3503
Epoch 4/10
1563/1563 [-----] - 23s 15ms/step - loss: 1.7895 - accuracy: 0.3505 - val_loss: 1.7763 - val_accuracy: 0.3530
Epoch 5/10
1563/1563 [-----] - 23s 15ms/step - loss: 1.7662 - accuracy: 0.3611 - val_loss: 1.7181 - val_accuracy: 0.3830
Epoch 6/10
1563/1563 [-----] - 23s 15ms/step - loss: 1.7506 - accuracy: 0.3663 - val_loss: 1.7250 - val_accuracy: 0.3832
Epoch 7/10
1563/1563 [-----] - 21s 13ms/step - loss: 1.7350 - accuracy: 0.3759 - val_loss: 1.6867 - val_accuracy: 0.3968
Epoch 8/10
1563/1563 [-----] - 21s 13ms/step - loss: 1.7189 - accuracy: 0.3812 - val_loss: 1.7239 - val_accuracy: 0.3777
Epoch 9/10
1563/1563 [-----] - 21s 13ms/step - loss: 1.7069 - accuracy: 0.3866 - val_loss: 1.7374 - val_accuracy: 0.3729
Epoch 10/10
1563/1563 [-----] - 21s 13ms/step - loss: 1.7019 - accuracy: 0.3882 - val_loss: 1.7405 - val_accuracy: 0.3734

```

```

313/313 [-----] - 3s 11ms/step - loss: 1.7405 - accuracy: 0.3734
Test accuracy: 0.3734

```

v.

Source	Magnitude (m = 0, n=0)		Magnitude (m = 2, n=3)		Magnitude (m = 4, n=7)	
Original Gabor Wavelet						

vi.



```
facial_expression_model_weights.h5 will be downloaded...
Downloading...
From: https://github.com/serengil/deepface\_models/releases/download/v1.0/facial\_expression\_model\_weights.h5
To: /root/.deepface/weights/facial_expression_model_weights.h5
100%|██████████| 5.98M/5.98M [00:00<00:00, 23.5MB/s]
{'emotion': {'angry': 7.147069602808642e-07, 'disgust': 1.7738306483383592e-12, 'fear': 7.72365460477431e-07, 'happy': 96.74765467643738, 'surprise': 1.147069602808642e-07, 'sadness': 1.147069602808642e-07}}
```

```
facial_expression_model_weights.h5 will be downloaded...
Downloading...
From: https://github.com/serengil/deepface\_models/releases/download/v1.0/facial\_expression\_model\_weights.h5
To: /root/.deepface/weights/facial_expression_model_weights.h5
100%|██████████| 5.98M/5.98M [00:00<00:00, 23.5MB/s]
{'emotion': {'angry': 7.147069602808642e-07, 'disgust': 1.7738306483383592e-12, 'fear': 7.72365460477431e-07, 'happy': 96.74765467643738, 'surprise': 1.147069602808642e-07, 'sadness': 1.147069602808642e-07}}
```

Conclusion: In this practical, I implemented deep learning concepts using DIGITS, TensorFlow, and PyTorch covered a wide spectrum of applications, including image classification, segmentation, object detection, transfer learning, face recognition, and emotion recognition